

# Job 2584773 - Gemma 2 9B IT Training (LoRA)

## Informazioni Generali

- **Job ID:** 2584773
- **Nome:** gemma\_te
- **Utente:** ediluzio
- **Status:** PENDING (in attesa priorità)
- **Partizione:** boost\_usr\_prod (GPU > 24GB)
- **Motivo Attesa:** (Priority) - Priorità inferiore rispetto ad altri job
- **Ambiente:** PyTorch 2.7.0+cu118 (riparato)

## Descrizione Tecnica

### Obiettivo del Job

Fine-tuning di Gemma 2 9B IT (Instruction Tuned) con LoRA per SVG captioning, continuando il training dal checkpoint esistente con ambiente PyTorch riparato.

### Architettura del Modello

#### Gemma 2 9B IT Base Model

Architecture: Transformer Decoder-Only (Gemma 2 family)  
- Layers: 42 transformer blocks  
- Hidden Size: 3584  
- Intermediate Size: 14336 (GeGLU activation)  
- Attention Heads: 16  
- Key-Value Heads: 8 (GQA – Grouped Query Attention)  
- Vocabulary Size: 256,000 tokens  
- Context Length: 8,192 tokens  
- Total Parameters: 9.24B  
- Instruction Tuned: Yes (IT variant)

#### Gemma 2 Innovations

```
# Key architectural improvements over Gemma 1
GEMMA2_FEATURES = {
    "sliding_window_attention": 4096, # Local attention pattern
    "soft_capping": 50.0,           # Attention logits capping
    "query_pre_attn_scalar": 3584**-0.5, # Query scaling
```

```

    "final_logit_softcapping": 30.0,      # Output logits capping
    "attn_logit_softcapping": 50.0,       # Attention logits capping
    "use_cache": True,                  # KV cache optimization
    "rope_theta": 10000.0               # RoPE base frequency
}

```

## LoRA Configuration (Optimized for Gemma 2)

```

LORA_CONFIG = {
    "r": 64,                      # Rank of adaptation
    "lora_alpha": 128,             # LoRA scaling parameter ( $\alpha/r = 2.0$ )
    "target_modules": [
        "q_proj", "k_proj", "v_proj", "o_proj",   # Attention projections
        "gate_proj", "up_proj", "down_proj"         # MLP projections
    ],
    "lora_dropout": 0.1,            # Dropout for LoRA layers
    "bias": "none",               # No bias in LoRA layers
    "task_type": "CAUSAL_LM",     # Causal language modeling
    "modules_to_save": ["embed_tokens", "lm_head"] # Full fine-tuning for embeddings
}

# Trainable Parameters Calculation
# Base model: 9.24B parameters (frozen)
# LoRA adapters: ~75M parameters (trainable)
# Embeddings: ~896M parameters (trainable if modules_to_save)
# Total trainable: ~75M (0.81%) or ~971M (10.5%) with embeddings

```

## Stack Tecnologico Riparato

### Environment Stack

- **PyTorch:** 2.7.0+cu118 (risolve CVE-2025-32434)
- **Transformers:** 4.52.3 (supporto completo Gemma 2)
- **PEFT:** 0.15.1 (LoRA implementation)
- **Accelerate:** 1.7.0 (distributed training)
- **BitsAndBytes:** 0.41.3 (quantization)
- **Flash Attention:** 2.x (memory efficient attention)

### Gemma 2 Specific Optimizations

```

# Attention optimizations
model_config.update({
    "use_flash_attention_2": True,      # Flash Attention 2.0
    "torch_dtype": torch.float16,       # Mixed precision
    "attnImplementation": "flash_attention_2"
})

```

```
# Memory optimizations for 9B model
torch.backends.cuda.matmul.allow_tf32 = True
torch.backends.cudnn.allow_tf32 = True
```

## Training Configuration

### Hyperparameters (Gemma 2 Optimized)

```
TRAINING_ARGS = {
    # Learning Rate & Optimization (Lower LR for larger model)
    "learning_rate": 1e-4,          # Reduced for 9B model stability
    "lr_scheduler_type": "cosine_with_restarts",
    "warmup_ratio": 0.05,          # Longer warmup for stability
    "weight_decay": 0.01,
    "adam_beta1": 0.9,
    "adam_beta2": 0.95,
    "adam_epsilon": 1e-8,

    # Training Dynamics
    "num_train_epochs": 2,         # Fewer epochs for larger model
    "per_device_train_batch_size": 1, # Smaller batch for 9B model
    "per_device_eval_batch_size": 2,
    "gradient_accumulation_steps": 16, # Higher accumulation
    "max_grad_norm": 0.5,          # Lower gradient clipping

    # Evaluation & Checkpointing
    "eval_strategy": "steps",
    "eval_steps": 10,
    "save_strategy": "steps",
    "save_steps": 10,
    "logging_steps": 1,
    "load_best_model_at_end": True,
    "metric_for_best_model": "eval_loss",

    # Early Stopping (More conservative)
    "early_stopping_patience": 10,
    "early_stopping_threshold": 0.0005,

    # Memory Optimization (Critical for 9B)
    "dataloader_pin_memory": True,
    "dataloader_num_workers": 2,   # Reduced for memory
    "remove_unused_columns": False,
    "fp16": True,
    "gradient_checkpointing": True,
    "optim": "adamw_torch_fused" # Fused optimizer
}
```

## Memory Management Strategy

```
# Advanced memory optimizations for 9B model
MEMORY_CONFIG = {
    "cpu_offload": False,           # Keep on GPU for speed
    "pin_memory": True,
    "empty_cache_steps": 10,        # Frequent cache clearing
    "max_memory_per_gpu": "28GB",   # Leave 4GB buffer
    "low_cpu_mem_usage": True,
    "torch_compile": True          # PyTorch 2.7 compilation
}
```

## Gemma 2 Specific Features

### Instruction Following Format

```
# Gemma 2 IT specific prompt format
def format_gemma_prompt(svg_xml, caption):
    prompt = "<bos><start_of_turn>user\n"
    prompt += "You are an expert at describing SVG images. "
    prompt += "Provide a detailed, accurate description of this SVG image:\n\n"
    prompt += f"{svg_xml}\n"
    prompt += "<end_of_turn>\n"
    prompt += "<start_of_turn>model\n"
    prompt += f"<{caption}>"
    prompt += "<end_of_turn><eos>"
    return prompt
```

```
# Special tokens for Gemma 2
SPECIAL_TOKENS = {
    "bos_token": "<bos>",
    "eos_token": "<eos>",
    "pad_token": "<pad>",
    "unk_token": "<unk>",
    "start_of_turn": "<start_of_turn>",
    "end_of_turn": "<end_of_turn>"
}
```

### Attention Pattern Optimization

```
# Sliding window attention for long sequences
def configure_sliding_window():
    model.config.sliding_window = 4096
    model.config.max_position_embeddings = 8192

    # Soft capping for numerical stability
    model.config.attn_logit_softcapping = 50.0
    model.config.final_logit_softcapping = 30.0
```

## Multi-GPU Training (Enhanced)

### Distributed Strategy for 9B Model

```
# Enhanced distributed setup for larger model
accelerate_config = {
    "compute_environment": "LOCAL_MACHINE",
    "distributed_type": "MULTI_GPU",
    "num_processes": 2,
    "gpu_ids": [0, 1],
    "mixed_precision": "fp16",
    "gradient_accumulation_steps": 16,
    "deepspeed_config": {
        "zero_stage": 2,           # ZeRO-2 for memory efficiency
        "offload_optimizer": False, # Keep optimizer on GPU
        "allgather_bucket_size": 5e8,
        "reduce_bucket_size": 5e8
    }
}
```

### Memory Requirements (9B Model)

Per GPU (>24GB richiesta):

- Model (9B params): ~18GB (fp16)
  - LoRA adapters: ~150MB (fp16)
  - Gradients: ~150MB
  - Optimizer states: ~300MB (AdamW)
  - Activations: ~6–8GB (batch\_size=1)
  - KV Cache: ~2GB
  - Buffer: ~2GB
- Total per GPU: ~28–30GB (richiede GPU >24GB)

## Checkpoint Management

### Resume from Existing Checkpoint

```
# Checkpoint path (may not exist yet)
checkpoint_path = "experiments/xml_direct_input/outputs/gemma2_9b_it_test3_multi_gpu/c

# Checkpoint verification
def verify_gemma_checkpoint():
    if os.path.exists(checkpoint_path):
        try:
            # Test compatibility with PyTorch 2.7.0
            state_dict = torch.load(f"{checkpoint_path}/pytorch_model.bin", weights_on
            print("✅ Gemma checkpoint compatible")
            return True
        except Exception as e:
```

```

        print(f"✗ Checkpoint error: {e}")
        return False
    else:
        print("⚠ Checkpoint not found, starting from base model")
        return False

# Training initialization
if verify_gemma_checkpoint():
    # Resume from checkpoint
    trainer.train(resume_from_checkpoint=checkpoint_path)
else:
    # Start fresh training
    trainer.train()

```

## Gemma 2 Specific Optimizations

### Flash Attention Integration

```

# Flash Attention 2.0 for memory efficiency
from flash_attn import flash_attn_func

class GemmaFlashAttention(nn.Module):
    def forward(self, query, key, value, attention_mask=None):
        # Flash attention with soft capping
        attn_output = flash_attn_func(
            query, key, value,
            dropout_p=0.0,
            softmax_scale=None,
            causal=True,
            window_size=(-1, -1),  # Full attention
            softcap=50.0            # Gemma 2 soft capping
        )
        return attn_output

```

### Quantization Support

```

# BitsAndBytes quantization for memory efficiency
from transformers import BitsAndBytesConfig

bnb_config = BitsAndBytesConfig(
    load_in_4bit=True,
    bnb_4bit_quant_type="nf4",
    bnb_4bit_compute_dtype=torch.float16,
    bnb_4bit_use_double_quant=True
)

# Load model with quantization
model = AutoModelForCausalLM.from_pretrained(
    "google/gemma-2-9b-it",

```

```
quantization_config=bnb_config,  
device_map="auto",  
torch_dtype=torch.float16  
)
```

## Performance Monitoring

### Gemma 2 Specific Metrics

```
# Additional metrics for Gemma 2  
MONITORING_METRICS = {  
    "attention_entropy": "Measure attention distribution",  
    "soft_capping_activation": "Monitor soft capping usage",  
    "sliding_window_efficiency": "Local attention effectiveness",  
    "instruction_following_score": "IT model specific metric",  
    "perplexity": "Language modeling quality"  
}
```

### Expected Performance

#### Training Characteristics:

- Slower convergence due to 9B **parameters**
- Better instruction following (IT variant)
- Higher memory requirements
- More stable training with soft capping

#### Performance Targets:

- BLEU-4: > 0.18 (higher than Llama due to IT)
- CLIP Score: > 0.28
- Instruction Following: > 0.85
- Training Time: 4-6 hours (larger model)

## Troubleshooting & Status

### Current Status

- **Environment:** PyTorch 2.7.0+cu118 ready
- **Model Support:** Transformers 4.52.3 supports Gemma 2
- **Memory Planning:** Optimized for >24GB GPUs
- **Queue Position:** Priority-based waiting

### Potential Issues & Solutions

1. **Memory Overflow:** Use gradient checkpointing + smaller batch size
2. **Slow Convergence:** Adjust learning rate schedule

**3. Instruction Format:** Ensure proper Gemma 2 IT formatting

**4. Attention Issues:** Monitor soft capping effectiveness

Questo job rappresenta il fine-tuning del modello più avanzato (Gemma 2 9B IT) con ottimizzazioni specifiche per l'architettura e capacità di instruction following.